

Multi-Class Classification				
Kaggle "Hate Speech and Offensive Language Dataset" dataset				
Model	Accuracy	Precision	Recall	F1-score
Distil BERT	0.91499	0.90835	0.91499	0.91073
Random Forest Classifier	0.87881	0.8671	0.87881	0.86959
CNN	0.87545	0.86423	0.87545	0.86709
LSTM	0.86644	0.85301	0.86644	0.85778
Random Forest Classifier (Baseline)	0.8308	0.82715	0.8308	0.82311
Binary Classification				
UC Berkley "Measuring Hate Speech" dataset from Huggingface				
Model	Accuracy	Precision	Recall	F1-score
Distil BERT	0.96633	0.96632	0.96633	0.96632
LSTM	0.9579	0.95782	0.9579	0.9578
CNN	0.95782	0.95782	0.95782	0.95766
Random Forest Classifier	0.94936	0.94977	0.94936	0.94896